


(DEEP) NEURAL NETWORKS



OUTLINE

- Introduction
- From neural networks to deep learning
- In practice:
 - Transfer learning
 - Standard libraries
- Limitations & future
- Focus: few-shot learning



OUTLINE

- Introduction
- From neural networks to deep learning
- In practice:
 - Transfer learning
 - Standard libraries
- Limitations & future
- Focus: few-shot learning



Introduction

└ What does deep learning mean?

https://en.wikipedia.org/wiki/Deep_learning




Deep learning (also known as deep structured learning or differential programming) is part of a broader family of **machine learning methods** based on **artificial neural networks** with representation learning. Learning can be supervised, semi-supervised or unsupervised.

[...]

A deep neural network (DNN) is an artificial neural network (ANN) with multiple layers between the input and output layers. The DNN finds the correct mathematical manipulation to turn the input into the output, whether it be a linear relationship or a non-linear relationship. The network moves through the layers calculating the probability of each output. [...] Each mathematical manipulation as such is considered a layer, and complex DNN have many layers, hence the name "deep" networks.

Introduction

└ Supervised, unsupervised, reinforcement




Introduction

└ A (very) growing field (1/2)

<https://medium.com/@dcharrezt/neurips-2019-stats-c91346d31c8f>

Statistics of acceptance rate NeurIPS




Introduction

└ A (very) growing field (2/2)

<https://www.kdnuggets.com/2018/10/recent-advances-deep-learning.html>

ICML stats for 2018



Introduction

└ Turing awards in 2018 (1/6)

<https://awards.acm.org/about/2018-turing>



Yoshua Bengio is a Professor at the **University of Montreal**, and the Scientific Director of both Mila (Quebec's Artificial Intelligence Institute) and IVADO (the Institute for Data Valorization). He is Co-director (with Yann LeCun) of CIFAR's Learning in Machines and Brains program. Bengio received a Bachelor's degree in electrical engineering, a Master's degree in computer science and a Doctoral degree in computer science from McGill University.

Bengio's honors include being named an Officer of the Order of Canada, Fellow of the Royal Society of Canada and the Marie-Victorin Prize. His work in founding and serving as Scientific Director of the Quebec Artificial Intelligence Institute (Mila) is also recognized as a major contribution to the field. Mila, an independent nonprofit organization, now counts 300 researchers and 35 faculty members among its ranks. It is the largest academic center for deep learning research in the world, and has helped put Montreal on the map as a vibrant AI ecosystem, with research labs from major companies as well as AI startups.

Introduction

└ Turing awards in 2018 (2/6)

<https://awards.acm.org/about/2018-turing>

Probabilistic models of sequences: In the 1990s, Bengio combined neural networks with probabilistic models of sequences, such as hidden Markov models. These ideas were incorporated into a system used by AT&T/NCR for reading handwritten checks, were considered a pinnacle of neural network research in the 1990s, and modern deep learning speech recognition systems are extending these concepts.

High-dimensional word embeddings and attention: In 2000, Bengio authored the landmark paper, “A Neural Probabilistic Language Model,” that introduced high-dimension word embeddings as a representation of word meaning. Bengio’s insights had a huge and lasting impact on natural language processing tasks including language translation, question answering, and visual question answering. His group also introduced a form of attention mechanism which led to breakthroughs in machine translation and form a key component of sequential processing with deep learning.

Generative adversarial networks: Since 2010, Bengio’s papers on generative deep learning, in particular the Generative Adversarial Networks (GANs) developed with Ian Goodfellow, have spawned a revolution in computer vision and computer graphics. In one fascinating application of this work, computers can actually create original images, reminiscent of the creativity that is considered a hallmark of human intelligence.

Introduction

└ Turing awards in 2018 (3/6)

<https://awards.acm.org/about/2018-turing>



Geoffrey Hinton is VP and Engineering Fellow of **Google**, Chief Scientific Adviser of The Vector Institute and a University Professor Emeritus at the University of Toronto. Hinton received a Bachelor's degree in experimental psychology from Cambridge University and a Doctoral degree in artificial intelligence from the University of Edinburgh. He was the founding Director of the Neural Computation and Adaptive Perception (later Learning in Machines and Brains) program at CIFAR.

Hinton's honors include Companion of the Order of Canada (Canada's highest honor), Fellow of the Royal Society (UK), foreign member of the National Academy of Engineering (US), the International Joint Conference on Artificial Intelligence (IJCAI) Award for Research Excellence, the NSERC Herzberg Gold medal, and the IEEE James Clerk Maxwell Gold medal. He was also selected by Wired magazine for "The Wired 100—2016's Most Influential People" and by Bloomberg for the 50 people who changed the landscape of global business in 2017.

Introduction

└ Turing awards in 2018 (3.5/6)

<https://awards.acm.org/about/2018-turing>



Geoffrey Hinton is VP and Engineering Fellow of **Google**, Chief Scientific Adviser of The Vector Institute and a University Professor Emeritus at the University of Toronto. Hinton received a Bachelor's degree in experimental psychology from Cambridge University and a Doctoral degree in artificial intelligence from the University of Edinburgh. He was the founding Director of the Neural Computation and Adaptive Perception (later Learning in Machines and Brains) program at CIFAR.

Hinton's honors include Companion of the Order of Canada (Canada's highest honor), Fellow of the Royal Society (UK), foreign member of the National Academy of Engineering (US), the International Joint Conference on Artificial Intelligence (IJCAI) Award for Research Excellence, the NSERC Herzberg Gold medal, and the IEEE James Clerk Maxwell Gold medal. He was also selected by Wired magazine for "The Wired 100—2016's Most Influential People" and by Bloomberg for the 50 people who changed the landscape of global business in 2017.

Introduction

└ Turing awards in 2018 (4/6)

<https://awards.acm.org/about/2018-turing>

Backpropagation: In a 1986 paper, “Learning Internal Representations by Error Propagation,” co-authored with David Rumelhart and Ronald Williams, Hinton demonstrated that the backpropagation algorithm allowed neural nets to discover their own internal representations of data, making it possible to use neural nets to solve problems that had previously been thought to be beyond their reach. The backpropagation algorithm is standard in most neural networks today.

Boltzmann Machines: In 1983, with Terrence Sejnowski, Hinton invented Boltzmann Machines, one of the first neural networks capable of learning internal representations in neurons that were not part of the input or output.

Improvements to convolutional neural networks: In 2012, with his students, Alex Krizhevsky and Ilya Sutskever, Hinton improved convolutional neural networks using rectified linear neurons and dropout regularization. In the prominent ImageNet competition, Hinton and his students almost halved the error rate for object recognition and reshaped the computer vision field.

Introduction

└ Turing awards in 2018 (5/6)

<https://awards.acm.org/about/2018-turing>



Yann LeCun is Silver Professor of the Courant Institute of Mathematical Sciences at New York University, and VP and Chief AI Scientist at **Facebook**. He received a Diplôme d'Ingénieur from the Ecole Supérieure d'Ingénieur en Electrotechnique et Electronique (ESIEE), and a PhD in computer science from Université Pierre et Marie Curie.

His honors include being a member of the US National Academy of Engineering; Doctorates Honoris Causa, from IPN Mexico and École Polytechnique Fédérale de Lausanne (EPFL); the Pender Award, University of Pennsylvania; the Holst Medal, Technical University of Eindhoven & Philips Labs; the Nokia-Bell Labs Shannon Luminary Award; the IEEE PAMI Distinguished Researcher Award; and the IEEE Neural Network Pioneer Award. He was also selected by Wired magazine for “The Wired 100—2016’s Most Influential People” and its “25 Geniuses Who are Creating the Future of Business.” LeCun was the founding director of the NYU Center of Data Science, and is a Co-director (with Yoshua Bengio) of CIFAR’s Learning in Machines and Brains program. LeCun is also a co-founder and former Member of the Board of the Partnership on AI, a group of companies and nonprofits studying the societal consequences of AI.

Introduction

└ Turing awards in 2018 (6/6)

<https://awards.acm.org/about/2018-turing>



Convolutional neural networks: In the 1980s, LeCun developed convolutional neural networks, a foundational principle in the field, which, among other advantages, have been essential in making deep learning more efficient. In the late 1980s, while working at the University of Toronto and Bell Labs, LeCun was the first to train a convolutional neural network system on images of handwritten digits. Today, convolutional neural networks are an industry standard in computer vision, as well as in speech recognition, speech synthesis, image synthesis, and natural language processing. They are used in a wide variety of applications, including autonomous driving, medical image analysis, voice-activated assistants, and information filtering.

Improving backpropagation algorithms: LeCun proposed an early version of the backpropagation algorithm (backprop), and gave a clean derivation of it based on variational principles. His work to speed up backpropagation algorithms included describing two simple methods to accelerate learning time.

Broadening the vision of neural networks: LeCun is also credited with developing a broader vision for neural networks as a computational model for a wide range of tasks, introducing in early work a number of concepts now fundamental in AI. For example, in the context of recognizing images, he studied how hierarchical feature representation can be learned in neural networks—a concept that is now routinely used in many recognition tasks. Together with Léon Bottou, he proposed the idea, used in every modern deep learning software, that learning systems can be built as complex networks of modules where backpropagation is performed through automatic differentiation. They also proposed deep learning architectures that can manipulate structured data, such as graphs.

Introduction

└ Nobel prizes in 2024 (chemistry)



John Michael Jumper, Demis Hassabis, David Baker

Introduction

└ Countless applications - Deep fakes

<https://www.youtube.com/watch?v=gLoI9hAX9dw>



Introduction

└ Countless applications - Defense

<https://www.clubic.com/drone/actualite-866669-essais-drones-prochaine-etape-reconnaissance-test-armee-americaine.html>

<https://talwaq1998.wordpress.com/2019/06/16/artificial-intelligence-in-the-military/>



Introduction

└ Countless applications - Risk estimation




<https://www.kaggle.com/c/prudential-life-insurance-assessment>
<https://www.hockani.com/blog/62/the-most-evil-insurance-companies/>



Introduction

└ Countless applications - Recognition

<https://www.intelligence-artificielle-school.com/actualite/deep-learning-intelligence-artificielle/>
<https://morioh.com/p/2ccee92141ef>



Introduction

└ Countless applications - Arts

Neural Rhapsody

maia (2019)

inspired by Wolfgang Mozart (1756-1791)

<https://towardsdatascience.com/neural-networks-for-music-generation-97c983b50204>

<https://www.lebigdata.fr/midjourney-concours-art>



Introduction

└ Countless applications - Natural language processing

<https://www.deepl.com>



GitHub
Copilot

A screenshot of the DeepL Translator website. The interface shows a translation from English to French. The English text on the left is: "The Planet That Once Used To Absorb Flesh In Order To Achieve Divinity And Immortality (Suffocated To The Flesh That It Desired...)" followed by a paragraph: "A silent forsaken planet. Its pores full of flesh. It once used to be alive. Now dead, condemned to loneliness. Lying in space, completely alone. Covered with tons and tons of piles of bone. The former master of the Universe and its stars. The former master died, its cover full of scars. Its core has gotten cool, the greed has shown fatality. No more absorbing, no consuming, it will never reach immortality". The French translation on the right is: "La planète qui autrefois absorbait la chair pour atteindre la divinité et l'immortalité (étouffée par la chair qu'elle désirait...)" followed by: "Une planète silencieuse et abandonnée. Ses pores sont pleins de chair. Elle était vivante autrefois. Maintenant morte, condamnée à la solitude. Allongée dans l'espace, complètement seule. Recouverte de tonnes et de tonnes de tas d'os. L'ancien maître de l'Univers et de ses étoiles. L'ancien maître est mort, sa couverture pleine de cicatrices. Son noyau s'est refroidi, la cupidité a montré la fatalité. Plus d'absorption, plus de consommation, elle n'atteindra jamais l'immortalité". The interface includes a navigation bar with "DeepL", "Traducteur", and "Linguee", a "Connexion" link, and a footer with "Ajouter un document", "542 / 5000", and icons for document, share, and download.

Introduction

└ Countless applications - Everyday life

https://actu.meilleurmobilite.com/google-assistant-alexa-siri-les-assistants-vocaux-au-coeur-dun-drole-de-test-de-piratage_200907

<https://gmail.googleblog.com/2015/07/the-mail-you-want-not-spam-you-dont.html>

<https://www.amazon.fr/SPAM-Classic-12-Ounce-Cans-Foods/dp/B001EQ5NHE>

<https://www.1min30.com/social-media-marketing/2148-2148>

<https://www.1min30.com/social-media-marketing/2148-2148>



f Machine Learning



Introduction



└ Countless applications - Self-driving cars

<https://www.eltis.org/in-brief/news/trust-and-safety-among-top-priorities-autonomous-vehicles>

<https://www.fia.com/fr/node/18761>

https://fr.wikipedia.org/wiki/Accident_de_la_route_en_France

Évolution du nombre de morts d'accidents de la route en France
de 1986 à 2016



Fu

22/10/2024

Introduction

└ Countless applications - Robotics

<https://emerj.com/ai-sector-overviews/machine-learning-in-robotics>





Fundamen

22/10/2024

Introduction

└ Countless applications - Space exploration

Morales et al. Deep learning for gravitational-wave data analysis: A resampling white-box approach. *Sensors* 21. 2021
<https://hyperight.com/how-nasa-uses-ai-and-machine-learning-for-space-exploration/>
<https://www.elliptika.com/fr/>




Fundamentals of M


22/10/2024

Introduction

└ Countless applications - Health



<https://qbrc.swmed.edu/projects/cnr/>
<https://alphafold.ebi.ac.uk/>
<https://www.wired.com/2014/07/everything-you-need-to-know-about-the-10-brain-myth-explained-in-60-seconds/>



Introduction

└ Why does deep learning work? - Universal approximator

Lu et al. The Expressive Power of Neural Networks: A View from the Width. *Advances in Neural Information Processing Systems* 30. 2017

Theorem 1 (Universal Approximation Theorem for Width-Bounded ReLU Networks). *For any Lebesgue-integrable function $f: \mathbb{R}^n \rightarrow \mathbb{R}$ and any $\epsilon > 0$, there exists a fully-connected ReLU network $\mathcal{A}$ with width $d_m \leq n + 4$, such that the function $F_{\mathcal{A}}$ represented by this network satisfies*


$$\int_{\mathbb{R}^n} |f(x) - F_{\mathcal{A}}(x)| dx < \epsilon. \quad (3)$$

In brief: Neural networks can be seen as universal function approximators.
i.e., they will succeed in most application cases

Introduction

└ Why does deep learning work? - Learns to abstract

<https://towardsdatascience.com/three-reasons-that-you-should-not-use-deep-learning-15bec517b622>



Funda

Feature visualization of convolutional net trained on ImageNet from [Zeiler & Fergus 2013]

22/10/2024


Introduction

└ Why does deep learning work? - The rise of GPUs



Fundamentals of Machine Learning

22/10/2024



OUTLINE

- Introduction
- From neural networks to deep learning
- In practice:
 - Transfer learning
 - Standard libraries
- Limitations & future
- Focus: few-shot learning




From neural networks to deep learning

└ Which ones do you know? (1/2)

From neural networks to deep learning

Which ones do you know? (2/2)



<https://www.asimovinstitute.org/author/fjodorvanveen>



From neural networks to deep learning

└ Taxonomy of neural networks - Perceptron (1/3)

<https://kamaleshjicse.files.wordpress.com/2018/11/pokedex-kamalesh.gif>





- Building block of everything that comes afterward
- Weights are trained by backpropagating the error
- Learns weights (coefficients) of a hyperplane to separate data points
- Handles linearly separable datasets only

From neural networks to deep learning

└ Taxonomy of neural networks - Perceptron (2/3)

<https://kamaleshjicse.files.wordpress.com/2018/11/pokedex-kamalesh.gif>




- Building block of everything that comes afterward
- Weights are trained by backpropagating the error
- Learns weights (coefficients) of a hyperplane to separate data points
- Handles linearly separable datasets only

From neural networks to deep learning


└ Taxonomy of neural networks - Perceptron (3/3)

<https://kamaleshjicse.files.wordpress.com/2018/11/pokedex-kamalesh.gif>

https://miro.medium.com/max/1800/1*c9WXUJdx00_fniK8h--k1A.gif



Example: Linear separability of points





- Building block of everything that comes afterward
- Weights are trained by backpropagating the error
- Learns weights (coefficients) of a hyperplane to separate data points
- Handles linearly separable datasets only

From neural networks to deep learning

└ Taxonomy of neural networks - MLP (1/3)

<https://kamaleshjicse.files.wordpress.com/2018/11/pokedex-kamalesh.gif>





- Direct extension of the perceptron
- Also named « feedforward networks »
- Vertical and/or horizontal duplication of neurons
- Handles non-linearly separable datasets

From neural networks to deep learning

└ Taxonomy of neural networks - MLP (2/3)

<https://kamaleshjicse.files.wordpress.com/2018/11/pokedex-kamalesh.gif>




- Direct extension of the perceptron
- Also named « feedforward networks »
- Vertical and/or horizontal duplication of neurons
- Handles non-linearly separable datasets


From neural networks to deep learning

└ Taxonomy of neural networks - MLP (3/3)

<https://kamaleshjicse.files.wordpress.com/2018/11/pokedex-kamalesh.gif>
https://miro.medium.com/max/935/1*XFSxzNm2LWbVxKsHxkFwiQ.png



Example: MNIST 2D projection





- Direct extension of the perceptron
- Also named « feedforward networks »
- Vertical and/or horizontal duplication of neurons
- Handles non-linearly separable datasets

From neural networks to deep learning

└ Taxonomy of neural networks - Autoencoder (1/3)

<https://kamaleshjicse.files.wordpress.com/2018/11/pokedex-kamalesh.gif>





- Output is trained to reproduce input
- Does not require labeled data → unsupervised
- Goes through a compressed representation of data
- Many extensions exist

From neural networks to deep learning

└ Taxonomy of neural networks - Autoencoder (2/3)

<https://kamaleshjicse.files.wordpress.com/2018/11/pokedex-kamalesh.gif>



- Output is trained to reproduce input
- Does not require labeled data → unsupervised
- Goes through a compressed representation of data
- Many extensions exist

From neural networks to deep learning


└ Taxonomy of neural networks - Autoencoder (3/3)

<https://kamaleshjicse.files.wordpress.com/2018/11/pokedex-kamalesh.gif>

Razavi et al. **Generating diverse high-fidelity images with vq-vae-2**. *Advances in Neural Information Processing Systems*. 2019



Example: Face generation





- Output is trained to reproduce input
- Does not require labeled data → unsupervised
- Goes through a compressed representation of data
- Many extensions exist

From neural networks to deep learning

└ Taxonomy of neural networks - RNN (1/3)

<https://kamaleshjicse.files.wordpress.com/2018/11/pokedex-kamalesh.gif>

<https://medium.com/@jon.froiland/advanced-use-of-recurrent-neural-networks-part-1-91ff1b36af25>





- Takes into account sequentiality of information (time series, texts...)
- Suffers a from gradient vanishing problem
- Base model for LSTMs or GRUs

From neural networks to deep learning

└ Taxonomy of neural networks - RNN (2/3)

<https://kamaleshjicse.files.wordpress.com/2018/11/pokedex-kamalesh.gif>

<https://medium.com/@jon.froiland/advanced-use-of-recurrent-neural-networks-part-1-91ff1b36af25>




- Takes into account sequentiality of information (time series, texts...)
- Suffers a from gradient vanishing problem
- Base model for LSTMs or GRUs

From neural networks to deep learning


└ Taxonomy of neural networks - RNN (3/3)

<https://kamaleshjicse.files.wordpress.com/2018/11/pokedex-kamalesh.gif>

<https://medium.com/@jon.froiland/advanced-use-of-recurrent-neural-networks-part-1-91ff1b36af25>



Example: Temperature forecasting





- Takes into account sequentiality of information (time series, texts...)
- Suffers a from gradient vanishing problem
- Base model for LSTMs or GRUs

From neural networks to deep learning

└ Taxonomy of neural networks - CNN (1/4)

<https://kamaleshjicse.files.wordpress.com/2018/11/pokedex-kamalesh.gif>





- Linear combination is replaced by convolution (captures locality of features)
- Some weights are shared (translation invariance)
- Ideal for data lying on Euclidean spaces
- Graph CNN extend this idea to irregular domains

From neural networks to deep learning

└ Taxonomy of neural networks - CNN (2/4)

<https://kamaleshjicse.files.wordpress.com/2018/11/pokedex-kamalesh.gif>




- Linear combination is replaced by convolution (captures locality of features)
- Some weights are shared (translation invariance)
- Ideal for data lying on Euclidean spaces
- Graph CNN extend this idea to irregular domains

From neural networks to deep learning

└ Taxonomy of neural networks - CNN (3/4)



<https://kamaleshjicse.files.wordpress.com/2018/11/pokedex-kamalesh.gif>

<https://towardsdatascience.com/intuitively-understanding-convolutions-for-deep-learning-1f6f42faee1>



3_0	3_1	2_2	1	0
0_2	0_2	1_0	3	1
3_0	1_1	2_2	2	3
2	0	0	2	2
2	0	0	0	1

12.0	12.0	17.0
10.0	17.0	19.0
9.0	6.0	14.0



- Linear combination is replaced by convolution (captures locality of features)
- Some weights are shared (translation invariance)
- Ideal for data lying on Euclidean spaces
- Graph CNN extend this idea to irregular domains


From neural networks to deep learning

└ Taxonomy of neural networks - CNN (4/4)

<https://kamaleshjicse.files.wordpress.com/2018/11/pokedex-kamalesh.gif>
<https://jonathan-hui.medium.com/image-segmentation-with-mask-r-cnn-ebe6d793272>



Example: Segmentation of images





- Linear combination is replaced by convolution (captures locality of features)
- Some weights are shared (translation invariance)
- Ideal for data lying on Euclidean spaces
- Graph CNN extend this idea to irregular domains

From neural networks to deep learning

└ Beyond those simple architectures

https://machinelearningmastery.com/wp-content/uploads/2021/08/attention_research_1.png


https://production-media.paperswithcode.com/methods/Screen_Shot_2021-01-26_at_9.43.31_PM_uI4jjMq.png



From neural networks to deep learning

└ Diffusion models

<https://theaisummer.com/diffusion-models/>



Fun

22/10/2024


From neural networks to deep learning

└ Foundation models

https://s46486.pcdn.co/wp-content/uploads/2023/03/YJCHsy_J6B4-00-02-42.png

A New Era of AI: Foundation Models

Step function improvements over legacy AI technologies




<https://github.com/brain-bzh/tutorial-foundation-models>

https://www.youtube.com/watch?v=0-pVVtLo4zY&list=PLfm_tlfDbIKvKCHs9nCXIgKFDBKpa6kdG


From neural networks to deep learning

└ Building blocks of a neural networks (1/6)




From neural networks to deep learning

└ Building blocks of a neural networks (2/6)




From neural networks to deep learning

└ Building blocks of a neural networks (3/6)




From neural networks to deep learning

└ Building blocks of a neural networks (4/6)



From neural networks to deep learning


└ Building blocks of a neural networks (5/6)



$$Y = f_3(W_3 \cdot f_2(W_2 \cdot f_1(W_1 \cdot X)))$$

From neural networks to deep learning

└ Building blocks of a neural networks (6/6)




$$Y = f_3(W_3 \cdot f_2(W_2 \cdot f_1(W_1 \cdot X)))$$

Derivable → Backpropagation


From neural networks to deep learning

└ Backpropagation (1/9)

<http://www.image-net.org>




Raw input
(or sometimes features)




From neural networks to deep learning

└ Backpropagation (2/9)

<http://www.image-net.org>





Raw input
(or sometimes features)



From neural networks to deep learning

└ Backpropagation (3/9)

<http://www.image-net.org>





(1) Forward pass

From neural networks to deep learning

└ Backpropagation (4/9)

<http://www.image-net.org>



$\tilde{y}_{beer}$ (expected)
(obtained)

$\tilde{y}_{wine}$ (expected)
(obtained)

(1) Forward pass

$$\epsilon = \text{loss}(\mathbf{Y}, \tilde{\mathbf{Y}})$$



(2) Compute error
using *loss function*

e.g., cross-entropy: $-\log(\mathbf{Y}^T \tilde{\mathbf{Y}})$

From neural networks to deep learning

└ Backpropagation (5/9)


<http://www.image-net.org>



From neural networks to deep learning

└ Backpropagation (6/9)



<http://www.image-net.org>



From neural networks to deep learning



└ Backpropagation (7/9)

<http://www.image-net.org>




From neural networks to deep learning

└ Backpropagation (8/9)



└ Backpropagation (9/9)


- Learning rate



From neural networks to deep learning

└ Training is performed in batches

<http://www.image-net.org>



From neural networks to deep learning


└ Batch normalization

Lofte et al. Batch Normalization: Accelerating Deep Network Training by Reducing Internal Covariate Shift. *International conference on machine learning*. 2015
<https://towardsdatascience.com/batch-normalization-in-3-levels-of-understanding-14c2da90a338#b93c>

Batch Normalization: Accelerating Deep Network Training by Reducing Internal Covariate Shift

Sergey Ioffe, Christian Szegedy

Training Deep Neural Networks is complicated by the fact that the distribution of each layer's inputs changes during training, as the parameters of the previous layers change. This slows down the training by requiring lower learning rates and careful parameter initialization, and makes it notoriously hard to train models with saturating nonlinearities. We refer to this phenomenon as internal covariate shift, and address the problem by normalizing layer inputs. Our method draws its strength from making normalization a part of the model architecture and performing the normalization for each training mini-batch. Batch Normalization allows us to use much higher learning rates and be less careful about initialization. It also acts as a regularizer, in some cases eliminating the need for Dropout. Applied to a state-of-the-art image classification model, Batch Normalization achieves the same accuracy with 14 times fewer training steps, and beats the original model by a significant margin. Using an ensemble of batch-normalized networks, we improve upon the best published result on ImageNet classification: reaching 4.9% top-5 validation error (and 4.8% test error), exceeding the accuracy of human raters.




From neural networks to deep learning

└ Regularizations during training


<https://medium.com/@amarbudhiraja/https-medium-com-amarbudhiraja-learning-less-to-learn-better-dropout-in-deep-machine-learning-74334da4bfc5>

<https://paperswithcode.com/method/weight-decay#>



(a) Standard Neural Net


Dropout




(b) After applying dropout.

Weight decay


Underfitting
(Excessive λ)



Appropriate weight decay
(Medium λ)



Overfitting
($\lambda \rightarrow 0$)



Regularization refers to training a model well so that it can generalize to data it hasn't seen before

Generalization usually prevents the model to converge too quickly


Fundamentals of Machine Learning

$$L_{new}(w) = L_{original}(w) + \lambda w^T w$$

From neural networks to deep learning

└ Training a neural network - Use a validation set

<http://www.cs.toronto.edu/~kriz/cifar.htm>




Fundamentals of Machine Learning

22/10/2024

From neural networks to deep learning

└ Training a neural network - Cross-validation & cross-testing


Credit: Amine El Ouahidi (Master report)



From neural networks to deep learning

└ Training a neural network - Beware of overfitting

https://upload.wikimedia.org/wikipedia/commons/1/1f/Overfitting_svg.svg



Train on 16000 samples, validate on 4000 samples

Epoch 1/200

16000/16000 [=====] - 3s 190us/step - loss: 0.1749 - val_loss: 0.1459

Epoch 00001: val_loss improved from inf to 0.14589, saving model to ...

Epoch 2/200

16000/16000 [=====] - 3s 173us/step - loss: 0.1282 - val_loss: 0.1159

Epoch 00002: val_loss improved from 0.14589 to 0.11591, saving model to ...

Epoch 3/200

16000/16000 [=====] - 3s 181us/step - loss: 0.0927 - val_loss: 0.0849

Epoch 00003: val_loss improved from 0.11591 to 0.08487, saving model to ...

Epoch 4/200

16000/16000 [=====] - 3s 181us/step - loss: 0.0608 - val_loss: 0.0594

Epoch 00004: val_loss improved from 0.08487 to 0.05943, saving model to ...

Epoch 5/200

16000/16000 [=====] - 3s 174us/step - loss: 0.0391 - val_loss: 0.0429

Epoch 00005: val_loss improved from 0.05943 to 0.04287, saving model to ...

Epoch 6/200

16000/16000 [=====] - 3s 177us/step - loss: 0.0258 - val_loss: 0.0324


Epoch 00006: val_loss improved from 0.04287 to 0.03238, saving model to ...

Epoch 7/200


15872/16000 [=====>.] - ETA: 0s - loss: 0.0175

From neural networks to deep learning


└ Training a neural network - Bias & variance



High bias
(underfit)




Well-trained



High variance
(overfit)

From neural networks to deep learning



└ Training a neural network - The problem of unbalanced classes (1/3)



From neural networks to deep learning

└ Training a neural network - The problem of unbalanced classes (2/3)



<https://towardsdatascience.com/breaking-the-curse-of-small-datasets-in-machine-learning-part-1-36f28b0c044d>



From neural networks to deep learning

└ Training a neural network - The problem of unbalanced classes (3/3)

<https://www.kdnuggets.com/2018/05/data-augmentation-deep-learning-limited-data.html>



Data augmentation
(rotating, scaling, brightness, color shift...)

From neural networks to deep learning

└ Training a neural network - Data augmentation is not only for class balancing


<https://neovision.fr/data-augmentation-solutions-manque-donnees>



From neural networks to deep learning

Self-supervised learning

Mangla et al. Charting the Right Manifold: Manifold Mixup for Few-shot Learning. IEEE/CVF Winter Conference on Applications of Computer Vision. 2020
<https://www.kdnuggets.com/2018/05/data-augmentation-deep-learning-limited-data.html>




Example 1: learn to quantify rotation

Adds *false* problems to solve

Trains more generic features

Can be done unsupervisedly

Example 2: learn to retrieve original image



OUTLINE

- Introduction
- From neural networks to deep learning
- In practice:
 - Transfer learning
 - Standard libraries
- Limitations & future
- Focus: few-shot learning



In practice – Transfer learning

Quiz (1/7)

You

Those big companies you all know

In practice – Transfer learning

└ Quiz (2/7)

You




Those big companies
you all know

In practice – Transfer learning

└ Quiz (3/7)

You





Those big companies
you all know


In practice – Transfer learning

Quiz (4/7)

You





Those big companies
you all know



In practice – Transfer learning

Quiz (5/7)

You




Those big companies
you all know


In practice – Transfer learning

Quiz (6/7)

You





Question 1: Who gets the best performance?



In practice – Transfer learning

Quiz (7/7)

You




Question 1: Who gets the best performance?
Question 2: Why not profit from what others have done?



In practice – Transfer learning

└ One more reason why you should consider transfer learning



<https://www.nature.com/articles/s42256-020-0219-9?proof=>
<https://www.forbes.com/sites/robtoews/2020/06/17/deep-learnings-climate-change-problem/>
<https://towardsdatascience.com/deep-learning-and-carbon-emissions-79723d5bc86e>

Common carbon footprint benchmarks

in lbs of CO2 equivalent

Roundtrip flight b/w NY and SF (1 passenger)

1,984

Human life (avg. 1 year)

11,023

American life (avg. 1 year)

36,156

US car including fuel (avg. 1 lifetime)

126,000

Transformer (213M parameters) w/ neural architecture search

626,155

Chart: MIT Technology Review • Source: Strubell et al. • Created with Datawrapper

In practice – Transfer learning

└ What is transfer learning? (1/2)


<https://towardsdatascience.com/deep-misconceptions-about-deep-learning-f26c41faceec>

Traditional ML


vs

Transfer Learning

- Isolated, single task learning:
 - Knowledge is not retained or accumulated. Learning is performed w.o. considering past learned knowledge in other tasks



- Learning of a new tasks relies on the previous learned tasks:
 - Learning process can be faster, more accurate and/or need less training data



In practice – Transfer learning

└ What is transfer learning? (2/2)

<https://towardsdatascience.com/a-comprehensive-hands-on-guide-to-transfer-learning-with-real-world-applications-in-deep-learning-212bf3b2f27a>


Transfer learning: idea

Instead of training a deep network from scratch for your task:

- Take a network trained on a different domain for a different **source task**
- Adapt it for your domain and your **target task**

Variations:

- Same domain, different task
- Different domain, same task



In practice – Transfer learning

└ Transfer learning in deep learning (1/3)

Example: same domain, different task


<https://hackernoon.com/traffic-sign-classification-6e7113d9c4d5>
<https://medium.com/zylapp/review-of-deep-learning-algorithms-for-image-classification-5fdbca4a05e2>
https://pytorch.org/tutorials/beginner/transfer_learning_tutorial.html

In practice – Transfer learning

└ Transfer learning in deep learning (2/3)

Example: same domain, different task

<https://hackernoon.com/traffic-sign-classification-6e7113d9c4d5>
<https://medium.com/zylapp/review-of-deep-learning-algorithms-for-image-classification-5fdbca4a05e2>
https://pytorch.org/tutorials/beginner/transfer_learning_tutorial.html



ResNet18 is a powerful image recognition network that has been trained on Imagenet 1000


Is of Machine Learning

In practice – Transfer learning

└ Transfer learning in deep learning (3/3)

Example: same domain, different task

<https://hackernoon.com/traffic-sign-classification-6e7113d9c4d5>
<https://medium.com/zylapp/review-of-deep-learning-algorithms-for-image-classification-5fdbca4a05e2>
https://pytorch.org/tutorials/beginner/transfer_learning_tutorial.html



ResNet18 is a powerful image recognition network that has been trained on Imagenet 1000

Objective: use its recognition power on our small dataset of bees and ants

predicted: bees



predicted: bees



predicted: bees



predicted: ants



predicted: bees




predicted: bees



Is of Machine Learning


In practice – Transfer learning

└ Exploiting a pre-trained model (1/7)




In practice – Transfer learning

└ Exploiting a pre-trained model (2/7)




In practice – Transfer learning

└ Exploiting a pre-trained model (3/7)




In practice – Transfer learning

└ Exploiting a pre-trained model (4/7)




In practice – Transfer learning

└ Exploiting a pre-trained model (5/7)




In practice – Transfer learning

└ Exploiting a pre-trained model (6/7)



In practice – Transfer learning

└ Exploiting a pre-trained model (7/7)



Now, we need to train the network a bit to make things work:

Solution 1: Freeze weights from the backbone and train the last layer

Solution 2 (fine-tuning): Train weights from both the backbone and the downstream

In practice – Transfer learning

└ Freezing layers or fine tuning?

<https://towardsdatascience.com/a-comprehensive-hands-on-guide-to-transfer-learning-with-real-world-applications-in-deep-learning-212bf3b2f27a>


Bottom n layers can be frozen or fine tuned.

- **Frozen:** not updated during backprop
- **Fine-tuned:** updated during backprop

Which to do depends on target task:

- **Freeze:** target task labels are scarce, and we want to avoid overfitting
- **Fine-tune:** target task labels are more plentiful

In general, we can set learning rates to be different for each layer to find a tradeoff between freezing and fine tuning



In practice – Transfer learning

└ Which backbone for which task? (1/2)

<https://modelzoo.co>

Categories

Computer Vision

Computer Vision: Object detection, boundary labelling, segmentation.

Natural Language Processing

Natural Language Processing (NLP).

Generative Models

Generative Models, such as Generative Adversarial Networks (GANs), Variational Autoencoders (VAE), and more.

Reinforcement Learning

Where an agent learn how to behave in a environment by performing actions and seeing the results.

Unsupervised Learning

Unsupervised learning is a type of machine learning algorithm to draw inferences from datasets consisting of input data without labels.


Audio and Speech

Models and code that perform audio processing, speech synthesis, and other audio related tasks.

In practice – Transfer learning

└ Which backbone for which task? (2/2)

<https://modelzoo.co>



Computer Vision

Computer Vision: Object detection, boundary labelling, segmentation.

Filter models...

OpenPose
★ 14173
OpenPose represents the first real-time multi-person system to jointly detect human body, hand, and facial keypoints (in total 130 keypoints) on single images.
Caffe
CV

Mask R-CNN
★ 13851
This is an implementation of Mask R-CNN on Python 3, Keras, and TensorFlow. The model generates bounding boxes and segmentation masks for each instance of an object in the image. It's based on Feature Pyramid Network (FPN) and a ResNet101 backbone.
Keras
CV

FastPhotoStyle
★ 9825
A Closed-form Solution to Photorealistic Image Stylization
PyTorch
CV

pytorch-CycleGAN-and-pix2pix
★ 9485
PyTorch implementation for both unpaired and paired image-to-image translation.
PyTorch
CV NLP Generative

Subscribe
We send out monthly emails showcasing the best or most notable models released each month. Sign up to receive updates!
your@email@domain.com
Subscribe

maskrcnn-benchmark
★ 6386
Fast, modular reference implementation of Instance Segmentation and Object Detection algorithms in PyTorch.
PyTorch
CV

OUTLINE

- Introduction
- From neural networks to deep learning
- **In practice:**
 - Transfer learning
 - **Standard libraries**
- Limitations & future
- Focus: few-shot learning



In practice – Standard libraries

A wide catalogue (1/4)

https://en.wikipedia.org/wiki/Comparison_of_deep-learning_software

Software	Creator	Initial Release	Software license ^[a]	Open source	Platform	Written in	Interface	OpenMP support	OpenCL support	CUDA support	Automatic differentiation ^[1]	Has pretrained models	Recurrent nets	Convolutional nets	RBM/DBNs	Parallel execution (multi node)	Actively Developed
BigDL	Jason Dai (Intel)	2016	Apache 2.0	Yes	Apache Spark	Scala	Scala, Python			No		Yes	Yes	Yes			
Caffe	Berkeley Vision and Learning Center	2013	BSD	Yes	Linux, macOS, Windows ^[2]	C++	Python, MATLAB, C++	Yes	Under development ^[3]	Yes	Yes	Yes ^[4]	Yes	Yes	No	?	No ^[5]
Chainer	Preferred Networks	2015	BSD	Yes	Linux, macOS	Python	Python	No	No	Yes	Yes	Yes	Yes	Yes	No	Yes	Yes
Deeplearning4j	Skymind engineering team; Deeplearning4j community; originally Adam Gibson	2014	Apache 2.0	Yes	Linux, macOS, Windows, Android (Cross-platform)	C++, Java	Java, Scala, Clojure, Python (Keras), Kotlin	Yes	No ^[6]	Yes ^{[7][8]}	Computational Graph	Yes ^[9]	Yes	Yes	Yes	Yes ^[10]	
Dlib	Davis King	2002	Boost Software License	Yes	Cross-Platform	C++	C++	Yes	No	Yes	Yes	Yes	No	Yes	Yes	Yes	
Intel Data Analytics Acceleration Library	Intel	2015	Apache License 2.0	Yes	Linux, macOS, Windows on Intel CPU ^[11]	C++, Python, Java	C++, Python, Java ^[11]	Yes	No	No	Yes	No		Yes		Yes	
Intel Math Kernel Library	Intel		Proprietary	No	Linux, macOS, Windows on Intel CPU ^[12]		C ^[13]	Yes ^[14]	No	No	Yes	No	Yes ^[15]	Yes ^[15]		No	
Keras	François Chollet	2015	MIT license	Yes	Linux, macOS, Windows	Python	Python, R	Only if using Theano as backend	Can use Theano, Tensorflow or PyTorch as backends	Yes	Yes	Yes ^[16]	Yes	Yes	No ^[17]	Yes ^[18]	Yes
MATLAB + Deep Learning Toolbox	MathWorks		Proprietary	No	Linux, macOS, Windows	C, C++, Java, MATLAB	MATLAB	No	No	Train with Parallel Computing Toolbox and generate CUDA code with GPU Coder ^[19]	Yes ^[20]	Yes ^{[21][22]}	Yes ^[21]	Yes ^[21]	Yes	With Parallel Computing Toolbox ^[23]	Yes
Microsoft Cognitive Toolkit (CNTK)	Microsoft Research	2016	MIT license ^[24]	Yes	Windows, Linux ^[25] (macOS via Docker on roadmap)	C++	Python (Keras), C++, Command line ^[26] BrainScript ^[27] (.NET on roadmap ^[28])	Yes ^[29]	No	Yes	Yes	Yes ^[30]	Yes ^[31]	Yes ^[31]	No ^[32]	Yes ^[33]	No ^[34]
Apache MXNet	Apache Software Foundation	2015	Apache 2.0	Yes	Linux, macOS, Windows, ^{[35][36]} AWS, Android, ^[37] iOS, JavaScript ^[38]	Small C++ core library	C++, Python, Julia, Matlab, JavaScript, Go, R, Scala, Perl, Clojure	Yes	On roadmap ^[39]	Yes	Yes ^[40]	Yes ^[41]	Yes	Yes	Yes	Yes ^[42]	Yes
Neural Designer	Arternics		Proprietary	No	Linux, macOS, Windows	C++	Graphical user interface	Yes	No	No	?	?	No	No	No	?	
OpenNN	Arternics	2003	GNU LGPL	Yes	Cross-platform	C++	C++	Yes	No	Yes	?	?	No	No	No	?	
PlaidML	Vertex AI, Intel	2017	AGPL	Yes	Linux, macOS, Windows	Python, C++, OpenCL	Python, C++	?	Some OpenCL ICDs are not recognized	No	Yes	Yes	Yes	Yes		Yes	Yes
PyTorch	Adam Paszke, Sam Gross, Soumith Chintala, Gregory Chanan (Facebook)	2016	BSD	Yes	Linux, macOS, Windows	Python, C, C++, CUDA	Python, C++	Yes	Via separately maintained package ^{[43][44][44]}	Yes	Yes	Yes	Yes	Yes		Yes	Yes
Apache SINGA	Apache Software Foundation	2015	Apache 2.0	Yes	Linux, macOS, Windows	C++	Python, C++, Java	No	Supported in V1.0	Yes	?	Yes	Yes	Yes	Yes	Yes	
TensorFlow	Google Brain	2015	Apache 2.0	Yes	Linux, macOS, Windows, ^[45] Android	C++, Python, CUDA	Python (Keras), C/C++, Java, Go, JavaScript, R, ^[46] Julia, Swift	No	On roadmap ^[47] but already with SYCL ^[48] support	Yes	Yes ^[49]	Yes ^[50]	Yes	Yes	Yes	Yes	Yes
Theano	Université de Montréal	2007	BSD	Yes	Cross-platform	Python	Python (Keras)	Yes	Under development ^[51]	Yes	Yes ^{[52][53]}	Through Lasagne's model zoo ^[54]	Yes	Yes	Yes	Yes ^[55]	No
Torch	Ronan Collobert, Koray Kavukcuoglu, Clement Faron	2002	BSD	Yes	Linux, macOS, Windows, ^[56] Android, ^[57] iOS	C, Lua	Lua, LuaJIT, ^[58] C, utility library for C++/OpenCL ^[59]	Yes	Third party implementations ^{[60][61]}	Yes ^{[62][63]}	Through Twitter's Autograd ^[64]	Yes ^[65]	Yes	Yes	Yes	Yes ^[66]	No
Wolfram Mathematica	Wolfram Research	1988	Proprietary	No	Windows, macOS, Linux, Cloud computing	C++, Wolfram Language, CUDA	Wolfram Language	Yes	No	Yes	Yes	Yes ^[66]	Yes	Yes	Yes	Yes ^[67]	Yes

In practice – Standard libraries

A wide catalogue (2/4)

https://en.wikipedia.org/wiki/Comparison_of_deep-learning_software

Software	Creator	Initial Release	Software license ^[a]	Open source	Platform	Written in	Interface	OpenMP support	OpenCL support	CUDA support	Automatic differentiation ^[1]	Has pretrained models	Recurrent nets	Convolutional nets	RBM/DBNs	Parallel execution (multi node)	Actively Developed
BigDL	Jason Dai (Intel)	2016	Apache 2.0	Yes	Apache Spark	Scala	Scala, Python			No		Yes	Yes	Yes			
Caffe	Berkeley Vision and Learning Center	2013	BSD	Yes	Linux, macOS, Windows ^[2]	C++	Python, MATLAB, C++	Yes	Under development ^[3]	Yes	Yes	Yes ^[4]	Yes	Yes	No	?	No ^[5]
Chainer	Preferred Networks	2015	BSD	Yes	Linux, macOS	Python	Python	No	No	Yes	Yes	Yes	Yes	Yes	No	Yes	Yes
Deeplearning4j	Skyrmind engineering team; Deeplearning4j community, origin Adam Gibson			Yes	Linux, macOS, Windows, Android (Cross-platform)	C++, Java	Java, Scala, Clojure, Python (Keras), Kotlin	Yes	No ^[6]	Yes ^{[7][8]}	Computational Graph	Yes ^[9]	Yes	Yes	Yes	Yes ^[10]	
Dlib	David E. King			Yes	Cross-Platform	C++	C++	Yes	No	Yes	Yes	Yes	No	Yes	Yes	Yes	
Intel Data Analytics Acceleration Library					macOS, Windows on Intel CPU ^[11]	C++, Python, Java	C++, Python, Java ^[11]	Yes	No	No	Yes	No		Yes		Yes	
Intel Math Kernel Library					macOS, Windows on Intel CPU ^[12]		C ^[13]	Yes ^[14]	No	No	Yes	No	Yes ^[15]	Yes ^[15]		No	
Keras	François Fleuret				Linux, macOS, Windows	Python	Python, R	Only if using Theano as backend	Can use Theano, Tensorflow or PyTorch as backends	Yes	Yes	Yes ^[16]	Yes	Yes	No ^[17]	Yes ^[18]	Yes
MATLAB + Deep Learning Toolbox	MathWorks		Proprietary	No	Linux, macOS, Windows	C, C++, Java, MATLAB	MATLAB	No	No	Train with Parallel Computing Toolbox and generate CUDA code with GPU Coder ^[19]	Yes ^[20]	Yes ^{[21][22]}	Yes ^[21]	Yes ^[21]	Yes	With Parallel Computing Toolbox ^[23]	Yes
Microsoft Cognitive Toolkit (CNTK)	Microsoft Research	2016	MIT license ^[24]	Yes	Windows, Linux ^[25] (macOS via Docker on roadmap)	C++	Python (Keras), C++, Command line ^[26] BrainScript ^[27] (.NET on roadmap ^[28])	Yes ^[29]	No	Yes	Yes	Yes ^[30]	Yes ^[31]	Yes ^[31]	No ^[32]	Yes ^[33]	No ^[34]
Apache MXNet	Apache Software Foundation	2015	Apache 2.0	Yes	Linux, macOS, Windows, ^{[35][36]} AWS, Android, ^[37] iOS, JavaScript ^[38]	Small C++ core library	C++, Python, Julia, Matlab, JavaScript, Go, R, Scala, Perl, Clojure	Yes	On roadmap ^[39]	Yes	Yes ^[40]	Yes ^[41]	Yes	Yes	Yes	Yes ^[42]	Yes
Neural Designer	Arternics		Proprietary	No	Linux, macOS, Windows	C++	Graphical user interface	Yes	No	No	?	?	No	No	No	?	
OpenNN	Arternics	2003	GNU LGPL	Yes	Cross-platform	C++	C++	Yes	No	Yes	?	?	No	No	No	?	
PlaidML	Vertex AI, Intel	2017	AGPL	Yes	Linux, macOS, Windows	Python, C++, OpenCL	Python, C++	?	Some OpenCL ICDs are not recognized	No	Yes	Yes	Yes	Yes		Yes	Yes
PyTorch	Adam Paszke, Sam Gross, Soumith Chintala, Gregory Chanan (Facebook)	2016	BSD	Yes	Linux, macOS, Windows	Python, C, C++, CUDA	Python, C++	Yes	Via separately maintained package ^{[43][44][44]}	Yes	Yes	Yes	Yes	Yes		Yes	Yes
Apache SINGA	Apache Software Foundation	2015	Apache 2.0	Yes	Linux, macOS, Windows	C++	Python, C++, Java	No	Supported in V1.0	Yes	?	Yes	Yes	Yes	Yes	Yes	
TensorFlow	Google Brain	2015	Apache 2.0	Yes	Linux, macOS, Windows, ^[45] Android	C++, Python, CUDA	Python (Keras), C/C++, Java, Go, JavaScript, R, ^[46] Julia, Swift	No	On roadmap ^[47] but already with SYCL ^[48] support	Yes	Yes ^[49]	Yes ^[50]	Yes	Yes	Yes	Yes	Yes
Theano	Université de Montréal	2007	BSD	Yes	Cross-platform	Python	Python (Keras)	Yes	Under development ^[51]	Yes	Yes ^{[52][53]}	Through Lasagne's model zoo ^[54]	Yes	Yes	Yes	Yes ^[55]	No
Torch	Ronan Collobert, Koray Kavukcuoglu, Clement Faron	2002	BSD	Yes	Linux, macOS, Windows, ^[56] Android, ^[57] iOS	C, Lua	Lua, LuaJIT, ^[58] C utility library for C++/OpenCL ^[59]	Yes	Third party implementations ^{[60][61]}	Yes ^{[62][63]}	Through Twitter's Autograd ^[64]	Yes ^[65]	Yes	Yes	Yes	Yes ^[66]	No
Wolfram Mathematica	Wolfram Research	1988	Proprietary	No	Windows, macOS, Linux, Cloud computing	C++, Wolfram Language, CUDA	Wolfram Language	Yes	No	Yes	Yes	Yes ^[66]	Yes	Yes	Yes	Yes ^[67]	Yes


TensorFlow
(Google Brain Team)
Not easy to learn
Bigger community
TensorBoard

In practice – Standard libraries

A wide catalogue (3/4)

https://en.wikipedia.org/wiki/Comparison_of_deep-learning_software

Software	Creator	Initial Release	Software license ^[a]	Open source	Platform	Written in	Interface	OpenMP support	OpenCL support	CUDA support	Automatic differentiation ^[1]	Has pretrained models	Recurrent nets	Convolutional nets	RBM/DBNs	Parallel execution (multi node)	Actively Developed
BigDL	Jason Dai (Intel)				Apache Spark	Scala	Scala, Python			No		Yes	Yes	Yes			
Caffe	Berkeley Vision and Learning Center				Linux, macOS, Windows ^[2]	C++	Python, MATLAB, C++	Yes	Under development ^[3]	Yes	Yes	Yes ^[4]	Yes	Yes	No	?	No ^[5]
Chainer	Preferred Networks				Linux, macOS	Python	Python	No	No	Yes	Yes	Yes	Yes	Yes	No	Yes	Yes
Deeplearning4j	Skymin DeepLearning				OS, Windows, Android, Cross-platform	C++, Java	Java, Scala, Clojure, Python (Keras), Kotlin	Yes	No ^[6]	Yes ^{[7][8]}	Computational Graph	Yes ^[9]	Yes	Yes	Yes	Yes ^[10]	
Dlib					Platform	C++	C++	Yes	No	Yes	Yes	Yes	No	Yes	Yes	Yes	
Intel Data Analytics Acceleration Library					Windows on Intel CPU ^[11]	C++, Python, Java	C++, Python, Java ^[11]	Yes	No	No	Yes	No		Yes		Yes	
Intel Math Kernel Library					OS, Windows on Intel CPU ^[12]		C ^[13]	Yes ^[14]	No	No	Yes	No	Yes ^[15]	Yes ^[15]		No	
Keras	François Fleuret				Linux, macOS, Windows	Python	Python, R	Only if using Theano as backend	Can use Theano, Tensorflow or PyTorch as backends	Yes	Yes	Yes ^[16]	Yes	Yes	No ^[17]	Yes ^[18]	Yes
MATLAB + Deep Learning Toolbox	MathWorks			No	Linux, macOS, Windows	C, C++, Java, MATLAB	MATLAB	No	No	Train with Parallel Computing Toolbox and generate CUDA code with GPU Coder ^[19]	Yes ^[20]	Yes ^{[21][22]}	Yes ^[21]	Yes ^[21]	Yes	With Parallel Computing Toolbox ^[23]	Yes
Microsoft Cognitive Toolkit (CNTK)	Microsoft Research	2016	MIT license ^[24]	Yes	Windows, Linux ^[25] (macOS via Docker on roadmap)	C++	Python (Keras), C++, Command line ^[26] BrainScript ^[27] (.NET on roadmap ^[28])	Yes ^[29]	No	Yes	Yes	Yes ^[30]	Yes ^[31]	Yes ^[31]	No ^[32]	Yes ^[33]	No ^[34]
Apache MXNet	Apache Software Foundation	2015	Apache 2.0	Yes	Linux, macOS, Windows, ^{[35][36]} AWS, Android, ^[37] iOS, JavaScript ^[38]	Small C++ core library	C++, Python, Julia, Matlab, JavaScript, Go, R, Scala, Perl, Clojure	Yes	On roadmap ^[39]	Yes	Yes ^[40]	Yes ^[41]	Yes	Yes	Yes	Yes ^[42]	Yes
Neural Designer	Arternics		Proprietary	No	Linux, macOS, Windows	C++	Graphical user interface	Yes	No	No	?	?	No	No	No	?	
OpenNN	Arternics	2003	GNU LGPL	Yes	Cross-platform	C++	C++	Yes	No	Yes	?	?	No	No	No	?	
PlaidML	Vertex AI, Intel	2017	AGPL	Yes	Linux, macOS, Windows	Python, C++, OpenCL	Python, C++	?	Some OpenCL ICDs are not recognized	No	Yes	Yes	Yes	Yes		Yes	Yes
PyTorch	Adam Paszke, Sam Gross, Soumith Chintala, Gregory Chanan (Facebook)	2016	BSD	Yes	Linux, macOS, Windows	Python, C, C++, CUDA	Python, C++	Yes	Via separately maintained package ^{[43][44][44]}	Yes	Yes	Yes	Yes	Yes		Yes	Yes
Apache SINGA	Apache Software Foundation	2015	Apache 2.0	Yes	Linux, macOS, Windows	C++	Python, C++, Java	No	Supported in V1.0	Yes	?	Yes	Yes	Yes	Yes	Yes	
TensorFlow	Google Brain	2015	Apache 2.0	Yes	Linux, macOS, Windows, ^[45] Android	C++, Python, CUDA	Python (Keras), C/C++, Java, Go, JavaScript, R, ^[46] Julia, Swift	No	On roadmap ^[47] but already with SYCL ^[48] support	Yes	Yes ^[49]	Yes ^[50]	Yes	Yes	Yes	Yes	Yes
Theano	Université de Montréal	2007	BSD	Yes	Cross-platform	Python	Python (Keras)	Yes	Under development ^[51]	Yes	Yes ^{[52][53]}	Through Lasagne's model zoo ^[54]	Yes	Yes	Yes	Yes ^[55]	No
Torch	Ronan Collobert, Koray Kavukcuoglu, Clement Faron	2002	BSD	Yes	Linux, macOS, Windows, ^[56] Android, ^[57] iOS	C, Lua	Lua, LuaJIT, ^[58] C utility library for C++/OpenCL ^[59]	Yes	Third party implementations ^{[60][61]}	Yes ^{[62][63]}	Through Twitter's Autograd ^[64]	Yes ^[65]	Yes	Yes	Yes	Yes ^[66]	No
Wolfram Mathematica	Wolfram Research	1988	Proprietary	No	Windows, macOS, Linux, Cloud computing	C++, Wolfram Language, CUDA	Wolfram Language	Yes	No	Yes	Yes	Yes ^[66]	Yes	Yes	Yes	Yes ^[67]	Yes

K Keras



TensorFlow

(Google Brain Team)

Not easy to learn
Bigger community
TensorBoard

In practice – Standard libraries

A wide catalogue (4/4)

https://en.wikipedia.org/wiki/Comparison_of_deep-learning_software

Software	Creator	Initial Release	Software license ^[a]	Open source	Platform	Written in	Interface	OpenMP support	OpenCL support	CUDA support	Automatic differentiation ^[1]	Has pretrained models	Recurrent nets	Convolutional nets	RBM/DBNs	Parallel execution (multi node)	Actively Developed
BigDL	Jason Dai (Intel)				Apache Spark	Scala	Scala, Python			No		Yes	Yes	Yes			
Caffe	Berkeley Vision and Learning Center				Linux, macOS, Windows ^[2]	C++	Python, MATLAB, C++	Yes	Under development ^[3]	Yes	Yes	Yes ^[4]	Yes	Yes	No	?	No ^[5]
Chainer	Preferred Networks				Linux, macOS	Python	Python	No	No	Yes	Yes	Yes	Yes	Yes	No	Yes	Yes
Deeplearning4j	Skymin DeepLearning				OS, Windows, Android (cross-platform)	C++, Java	Java, Scala, Clojure, Python (Keras), Kotlin	Yes	No ^[6]	Yes ^{[7][8]}	Computational Graph	Yes ^[9]	Yes	Yes	Yes	Yes ^[10]	
Dlib					Platform	C++	C++	Yes	No	Yes	Yes	Yes	No	Yes	Yes	Yes	
Intel Data Analytics Acceleration Library					Windows on Intel CPU ^[11]	C++, Python, Java	C++, Python, Java ^[11]	Yes	No	No	Yes	No		Yes		Yes	
Intel Math Kernel Library					OS, Windows on Intel CPU ^[12]		C ^[13]	Yes ^[14]	No	No	Yes	No	Yes ^[15]	Yes ^[15]		No	
Keras	François Fleuret				Linux, macOS, Windows	Python	Python, R	Only if using Theano as backend	Can use Theano, Tensorflow or PyTorch as backends	Yes	Yes	Yes ^[16]	Yes	Yes	No ^[17]	Yes ^[18]	Yes
MATLAB + Deep Learning Toolbox	MathWorks			No	Linux, macOS, Windows	C, C++, Java, MATLAB	MATLAB	No	No	Train with Parallel Computing Toolbox and generate CUDA code with GPU Coder ^[19]	Yes ^[20]	Yes ^{[21][22]}	Yes ^[21]	Yes ^[21]	Yes	With Parallel Computing Toolbox ^[23]	Yes
Microsoft Cognitive Toolkit (CNTK)	Microsoft Research	2016	MIT license ^[24]	Yes	Windows, Linux ^[25] (macOS via Docker)	C++	Python (Keras), C++, Command line ^[26] BrainScript ^[27] (.NET on roadmap ^[28])	Yes ^[29]	No	Yes	Yes	Yes ^[30]	Yes ^[31]	Yes ^[31]	No ^[32]	Yes ^[33]	No ^[34]
Apache MXNet	Apache Software Foundation	2015	Apache 2.0	Yes	Linux, macOS, Windows	C++ core library	C++, Python, Julia, Matlab, JavaScript, Go, R, Scala, Perl, Clojure	Yes	On roadmap ^[38]	Yes	Yes ^[40]	Yes ^[41]	Yes	Yes	Yes	Yes ^[42]	Yes
Neural Designer	Arternics						Graphical user interface	Yes	No	No	?	?	No	No	No	?	
OpenNN	Arternics	2003					C++	Yes	No	Yes	?	?	No	No	No	?	
PlaidML	Vertex AI, Intel	2017				C++, C, C++	Python, C++	?	Some OpenCL ICDs are not recognized	No	Yes	Yes	Yes	Yes		Yes	Yes
PyTorch	Adam Paszke, Sam Gross, Soumith Chintala, Gregory Chanan (Facebook)	2016				C, C++, CUDA	Python, C++	Yes	Via separately maintained package ^{[43][44][44]}	Yes	Yes	Yes	Yes	Yes		Yes	Yes
Apache SINGA	Apache Software Foundation	2015	Apache 2.0	Yes	Linux, macOS, Windows	C++	Python, C++, Java	No	Supported in V1.0	Yes	?	Yes	Yes	Yes	Yes	Yes	
TensorFlow	Google Brain	2015	Apache 2.0	Yes	Android	C++, Python, CUDA	Python (Keras), C/C++, Java, Go, JavaScript, R ^[46] Julia, Swift	No	On roadmap ^[47] but already with SYCL ^[48] support	Yes	Yes ^[49]	Yes ^[50]	Yes	Yes	Yes	Yes	Yes
Theano	Université de Montréal	2007	BSD	Yes	Cross-platform	Python	Python (Keras)	Yes	Under development ^[51]	Yes	Yes ^{[52][53]}	Through Lasagne's model zoo ^[54]	Yes	Yes	Yes	Yes ^[55]	No
Torch	Ronan Collobert, Koray Kavukcuoglu, Clement Faron	2002	BSD	Yes	Linux, macOS, Windows ^[56] Android ^[57] iOS	C, Lua	Lua, LuaJIT ^[58] C, utility library for C++/OpenCL ^[59]	Yes	Third party implementations ^{[60][61]}	Yes ^{[62][63]}	Through Twitter's Autograd ^[64]	Yes ^[65]	Yes	Yes	Yes	Yes ^[66]	No
Wolfram Mathematica	Wolfram Research	1988	Proprietary	No	Windows, macOS, Linux, Cloud computing	C++, Wolfram Language, CUDA	Wolfram Language	Yes	No	Yes	Yes	Yes ^[66]	Yes	Yes	Yes	Yes ^[67]	Yes

K Keras



TensorFlow

(Google Brain Team)

Not easy to learn
Bigger community
TensorBoard

PYTORCH

(Facebook AI Research)

Production-ready
Strong community
Faster than Keras

In practice – Standard libraries

└ PyTorch vs. TensorFlow (1/3)

<https://www.youtube.com/watch?v=Mn1AACvK-YI>



Fundam


22/10/2024

In practice – Standard libraries

└ PyTorch vs. TensorFlow (2/3)

<https://www.youtube.com/watch?v=Mn1AACvK-YI>

Conclusion – My Opinion



PyTorch

- PyTorch is better for **rapid prototyping** in Research.
- Good start for **hobbyists** and **small-scale** projects
- Better **utilization** and **optimization** for use of **GPUs**

TensorFlow

- TensorFlow is better for **large-scale** deployments
- Good for **cross-platform** and **embedded** deployment
- Bigger** community base and **many** users

65%

35%

edureka!

AI & Deep Learning Training

www.edureka.co/ai-deep-learning


SUBSCRIBE

In practice – Standard libraries

└ PyTorch vs. TensorFlow (3/3)

<https://www.youtube.com/watch?v=Mn1AACvK-Y>

Today's
practical session



Later, when you're
a data scientist

... if you
really need it ...

In practice – Standard libraries

└ A quick example in PyTorch - Step 1: prepare the dataloaders

https://pytorch.org/tutorials/beginner/blitz/cifar10_tutorial.html

```
transform = transforms.Compose(
    [transforms.ToTensor(),
     transforms.Normalize((0.5, 0.5, 0.5), (0.5, 0.5, 0.5))])

batch_size = 4

trainset = torchvision.datasets.CIFAR10(root='./data', train=True,
                                         download=True, transform=transform)
trainloader = torch.utils.data.DataLoader(trainset, batch_size=batch_size,
                                           shuffle=True, num_workers=2)

testset = torchvision.datasets.CIFAR10(root='./data', train=False,
                                         download=True, transform=transform)
testloader = torch.utils.data.DataLoader(testset, batch_size=batch_size,
                                          shuffle=False, num_workers=2)

classes = ('plane', 'car', 'bird', 'cat',
           'deer', 'dog', 'frog', 'horse', 'ship', 'truck')
```

In practice – Standard libraries

└ A quick example in PyTorch - Step 2: Define the model

https://pytorch.org/tutorials/beginner/blitz/cifar10_tutorial.html

```
class Net(nn.Module):
    def __init__(self):
        super().__init__()
        self.conv1 = nn.Conv2d(3, 6, 5)
        self.pool = nn.MaxPool2d(2, 2)
        self.conv2 = nn.Conv2d(6, 16, 5)
        self.fc1 = nn.Linear(16 * 5 * 5, 120)
        self.fc2 = nn.Linear(120, 84)
        self.fc3 = nn.Linear(84, 10)

    def forward(self, x):
        x = self.pool(F.relu(self.conv1(x)))
        x = self.pool(F.relu(self.conv2(x)))
        x = torch.flatten(x, 1) # flatten all dimensions except batch
        x = F.relu(self.fc1(x))
        x = F.relu(self.fc2(x))
        x = self.fc3(x)
        return x

net = Net()
```

In practice – Standard libraries

└ A quick example in PyTorch - Step 3: Choose the loss function, optimizer, scheduler, etc.

https://pytorch.org/tutorials/beginner/blitz/cifar10_tutorial.html

```
import torch.optim as optim

criterion = nn.CrossEntropyLoss()
optimizer = optim.SGD(net.parameters(), lr=0.001, momentum=0.9)
```


In practice – Standard libraries

└ A quick example in PyTorch - Step 4: Train the model

[tutorials/beginner/blitz/cifar10_tutorial.html](https://pytorch.org/tutorials/beginner/blitz/cifar10_tutorial.html)

```
for epoch in range(2): # loop over the dataset multiple times

    running_loss = 0.0
    for i, data in enumerate(trainloader, 0):
        # get the inputs; data is a list of [inputs, labels]
        inputs, labels = data

        # zero the parameter gradients
        optimizer.zero_grad()

        # forward + backward + optimize
        outputs = net(inputs)
        loss = criterion(outputs, labels)
        loss.backward()
        optimizer.step()

        # print statistics
        running_loss += loss.item()
        if i % 2000 == 1999: # print every 2000 mini-batches
            print(f'[epoch + 1], {i + 1:5d} loss: {running_loss / 2000:.3f}')
            running_loss = 0.0

print('Finished Training')
```

You can evaluate the model on the validation dataset here

In practice – Standard libraries

└ A quick example in PyTorch - Step 5: Test the model

https://pytorch.org/tutorials/beginner/blitz/cifar10_tutorial.html

```
correct = 0
total = 0
# since we're not training, we don't need to calculate the gradients for our outputs
with torch.no_grad():
    for data in testloader:
        images, labels = data
        # calculate outputs by running images through the network
        outputs = net(images)
        # the class with the highest energy is what we choose as prediction
        _, predicted = torch.max(outputs.data, 1)
        total += labels.size(0)
        correct += (predicted == labels).sum().item()

print(f'Accuracy of the network on the 10000 test images: {100 * correct // total} %')
```

OUTLINE


- Introduction
- From neural networks to deep learning
- In practice:
 - Transfer learning
 - Standard libraries
- **Limitations & future**
- Focus: few-shot learning



Limitations & future

- └ The limits of deep learning - It (generally) requires a lot of data

<https://towardsdatascience.com/deep-misconceptions-about-deep-learning-f26c41faceec>



Limitations & future

- └ The limits of deep learning - It requires similarities in data distributions

<https://www.theverge.com/2016/4/7/11383272/google-self-driving-arizona-desert-conditions>

<https://www.pinterest.fr/pin/476255729316690445/>





Limitations & future

- └ The limits of deep learning - It is not easy to train

<https://xkcd.com/1838>

<https://towardsdatascience.com/three-reasons-that-you-should-not-use-deep-learning-15bec517b622>





Limitations & future

└ The limits of deep learning - Adversarial examples

<https://openai.com/blog/adversarial-example-research>

<https://medium.com/syncedreview/natural-adversarial-examples-slash-image-classification-accuracy-by-90-702f381acbb8>

https://nicholas.carlini.com/code/audio_adversarial_examples



▶ 0:03 / 0:03 — 🔊 ⋮

"okay google browse to evil dot com"

▶ 0:03 / 0:03 — 🔊 ⋮

"without the dataset the article is useless"


Fundamentals of Machine Learning

22/10/2024

Limitations & future

└ The limits of deep learning - The challenge of interpretability


<https://towardsdatascience.com/interpretability-vs-accuracy-the-friction-that-defines-deep-learning-dae16c84db5c>



Limitations & future

- └ The limits of deep learning - Something obscure for society, exploited by quacks

<https://www.youtube.com/watch?app=desktop&v=miIy2gb-Lgc&pts.com/laurent-alexandre-premunir-contre-dangers-de-intelligence-artificielle/>



Limitations & future

└ The limits of deep learning - A strong environmental impact

<https://news.mit.edu/2020/shrinking-deep-learning-carbon-footprint-0807>
<https://arxiv.org/pdf/1906.02243.pdf>

In June, OpenAI unveiled the largest language model in the world, a text-generating tool called GPT-3 that can write creative fiction, translate legalese into plain English, and answer obscure trivia questions. It's the latest feat of intelligence achieved by deep learning, a machine learning method patterned after the way neurons in the brain process and store information.

But it came at a hefty price: at least \$4.6 million and 355 years in computing time, assuming the model was trained on a standard neural network chip, or GPU. The model's colossal size — 1,000 times larger than a typical language model — is the main factor in its high cost.


Consumption	CO ₂ e (lbs)
Air travel, 1 passenger, NY↔SF	1984
Human life, avg, 1 year	11,023
American life, avg, 1 year	36,156
Car, avg incl. fuel, 1 lifetime	126,000
Training one model (GPU)	
NLP pipeline (parsing, SRL)	39
w/ tuning & experimentation	78,468
Transformer (big)	192
w/ neural architecture search	626,155

Table 1: Estimated CO₂ emissions from training common NLP models, compared to familiar consumption.¹

Limitations & future

└ Research in deep learning - Few-shot learning


<https://neptune.ai/blog/understanding-few-shot-learning-in-computer-vision>



Limitations & future

└ Research in deep learning - Active learning


<https://blogs.nvidia.com/blog/2020/01/16/what-is-active-learning>



Limitations & future

└ Research in deep learning - Multimodal deep learning


<https://arxiv.org/pdf/2105.11087.pdf>




Limitations & future

└ Research in deep learning - Geometric deep learning


<https://flawnsontong.medium.com/what-is-geometric-deep-learning-b2adb662d91d>




Molecules



Trees



Networks





Manifolds

Limitations & future

└ Research in deep learning - Compression of models


<https://medium.com/gsi-technology/an-overview-of-model-compression-techniques-for-deep-learning-in-space-3fd8d4ce84e5>




Limitations & future

└ Research in deep learning - Curriculum learning

<https://arxiv.org/pdf/2101.10382.pdf>



a General framework for data-level curriculum learning.




b General framework for model-level curriculum learning.

Limitations & future


└ Research in deep learning - Continual learning

<https://medium.com/continual-ai/towards-adaptive-ai-with-continual-learning-f493fd0d698>

Static ML



Adaptive ML



Limitations & future

└ Research in deep learning - Interpretability / explicability

<https://arxiv.org/pdf/2103.10689.pdf>




(a) Image



(b)
Human Label



(c) LIME



(d) GradCAM



(e) SmoothGrad



(f) Image



(g)
Human Label



(h) LIME



(i) GradCAM



(j) SmoothGrad

Limitations & future

└ Research in deep learning - Robustness

https://infoscience.epfl.ch/record/229872/files/spm_preprint.pdf

<https://upcommons.upc.edu/bitstream/handle/2117/188799/Givaki%20et%20al.pdf>





FIGURE 6. r_{adv}^* denotes the adversarial perturbation of x (with $p = 2$). Note that r_{adv}^* is orthogonal to the decision boundary $\mathcal{B}$ and $\|r_{\text{adv}}^*\|_2 = \text{dist}(x, \mathcal{B})$.




(c) (Activation Function, FPGA)= (Tanh,VC707)

Limitations & future

└ Research in deep learning - Etc.

https://www.google.com



The image shows a Google search interface. The search bar contains the text "research topics in deep learning". Below the search bar, there are tabs for "Tous", "Images", "Vidéos", "Actualités", "Livres", and "Plus". The "Tous" tab is selected. Below the tabs, it says "Environ 372 000 000 résultats (0,64 secondes)". The search results are titled "Articles universitaires correspondant aux termes **research topics in deep learning**". There are three results listed: "Deep learning - Goodfellow - Cité 43293 fois", "Deep learning in bioinformatics - Min - Cité 1286 fois", and "Deep learning - Bengio - Cité 1447 fois". Below these results, there is a section titled "This article features the top 10 research and thesis topics for ML projects for students to try in 2022". This section contains a list of 10 topics: "Text Mining and Text Classification. ...", "Image-Based Applications. ...", "Machine Vision. ...", "Clustering. ...", "Optimization. ...", "Voice Classification. ...", "Sentiment Analysis. ...", and "Recommendation Framework Project.".

Google

research topics in deep learning

Tous Images Vidéos Actualités Livres Plus Outils

Environ 372 000 000 résultats (0,64 secondes)

Articles universitaires correspondant aux termes **research topics in deep learning**

Deep learning - Goodfellow - Cité 43293 fois

Deep learning in bioinformatics - Min - Cité 1286 fois

Deep learning - Bengio - Cité 1447 fois

This article features the top 10 research and thesis topics for ML projects for students to try in 2022

- Text Mining and Text Classification. ...
- Image-Based Applications. ...
- Machine Vision. ...
- Clustering. ...
- Optimization. ...
- Voice Classification. ...
- Sentiment Analysis. ...
- Recommendation Framework Project.

Plus... • 7 janv. 2022

Fundamentals of

22/10/2024

OUTLINE

- Introduction
- From neural networks to deep learning
- In practice:
 - Transfer learning
 - Standard libraries
- Limitations & future
- Focus: few-shot learning



Few-shot learning

└ Motivating example


Tenenbaum, J. B., Kemp, C., Griffiths, T. L., & Goodman, N. D. (2011). How to grow a mind: Statistics, structure, and abstraction. *science*, 331(6022), 1279-1285



Few-shot learning

└ Recall: supervised algorithms require a lot of data

<https://towardsdatascience.com/deep-misconceptions-about-deep-learning-f26c41faceec>




Few-shot learning

└ The difference with transfer learning (1/2)

<https://hackernoon.com/traffic-sign-classification-6e7113d9c4d5>

<https://medium.com/zylapp/review-of-deep-learning-algorithms-for-image-classification-5fdbca4a05e2>

https://pytorch.org/tutorials/beginner/transfer_learning_tutorial.html



ResNet18 is a powerful image recognition network that has been trained on Imagenet 1000

Objective: use its recognition power on our small dataset of bees and ants

predicted: bees



predicted: bees



predicted: bees



predicted: ants



predicted: bees



predicted: bees



Is of Machine Learning


Few-shot learning

└ The difference with transfer learning (2/2)

<https://hackernoon.com/traffic-sign-classification-6e7113d9c4d5>

<https://medium.com/zylapp/review-of-deep-learning-algorithms-for-image-classification-5fdbca4a05e2>

https://pytorch.org/tutorials/beginner/transfer_learning_tutorial.html



ResNet18 is a powerful image recognition network that has been trained on Imagenet 1000

Objective: use its recognition power on our small dataset of bees and ants

predicted: bees



predicted: bees



predicted: bees



predicted: ants



predicted: bees



predicted: bees



Typically 1-5 labeled samples per class

Is of Machine Learning




Few-shot learning

└ When does it occur?

<https://www.semanticscholar.org/paper/Navigating-a-smart-wheelchair-with-a-brain-computer-Mandel-L%C3%BCth/8b6c1b061bbc3c54186a06bb9cc40a0367eb9d44>

<https://www.nature.com/articles/s43018-020-00169-2>


Bonardi et al. **Learning One-Shot Imitation from Humans without Humans.** *IEEE Robotics and Automation Letters*, 5, 2020



ng


Few-shot learning

- └ Few-shot vocabulary (1/8)




Few-shot learning

└ Few-shot vocabulary (2/8)




D_{base}



D_{novel}


Few-shot learning

- └ Few-shot vocabulary (3/8)




Few-shot learning

└ Few-shot vocabulary (4/8)



D_{base}




D_{novel}


➤ N ways:
Unseen classes
Here, 5

Few-shot learning

└ Few-shot vocabulary (5/8)



D_{base}




D_{novel}


- N ways:
Unseen classes
Here, 5
- K shots:
Labeled samples per class
Here, 1

Few-shot learning

└ Few-shot vocabulary (6/8)



D_{base}




D_{novel}

Support set of NK samples


- N ways:
Unseen classes
Here, 5
- K shots:
Labeled samples per class
Here, 1

Few-shot learning

└ Few-shot vocabulary (7/8)



D_{base}





D_{novel}

Support set of NK samples

- N ways:
Unseen classes
Here, 5
 - K shots:
Labeled samples per class
Here, 1
- Queries:
Samples to classify
From selected ways

Few-shot learning

└ Few-shot vocabulary (8/8)





Support set of NK samples

- N ways:
Unseen classes
Here, 5
- K shots:
Labeled samples per class
Here, 1
- Queries:
Samples to classify
From selected ways



Few-shot learning

- Inductive few-shot learning



Few-shot learning


- └ Transductive few-shot learning



Few-shot learning

└ The classical pipeline


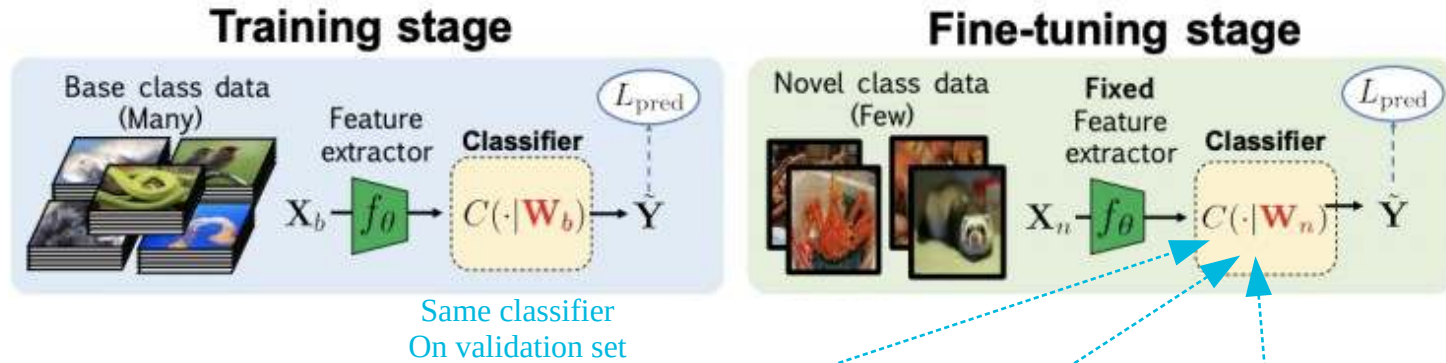
Credit: Raphael Lafargue



Few-shot learning

└ Examples of classifiers


<https://opencv.org/understanding-transductive-few-shot-learning>
<https://spin.atomicobject.com/2015/05/26/mean-shift-clustering/>



Nearest class mean



Label propagation




Mean shift

Few-shot learning

The EASY approach - Description

<https://arxiv.org/pdf/2201.09699.pdf>



Few-shot learning

The EASY approach - Results

<https://arxiv.org/pdf/2201.09699.pdf>

TABLE I

1-SHOT AND 5-SHOT ACCURACY OF STATE-OF-THE-ART METHODS AND PROPOSED SOLUTION ON **MINIIMAGENET** IN **INDUCTIVE** SETTING.

Method	1-shot	5-shot	
$\leq 12M$	SimpleShot [29]	62.85 ± 0.20	80.02 ± 0.14
	Baseline++ [30]	53.97 ± 0.79	75.90 ± 0.61
	TADAM [35]	58.50 ± 0.30	76.70 ± 0.30
	ProtoNet [10]	60.37 ± 0.83	78.02 ± 0.57
	R2-D2 (+ens) [20]	64.79 ± 0.45	81.08 ± 0.32
	FEAT [36]	66.78	82.05
	CNL [37]	67.96 ± 0.98	83.36 ± 0.51
	MERL [38]	67.40 ± 0.43	83.40 ± 0.28
	Deep EMD v2 [13]	68.77 ± 0.29	84.13 ± 0.53
	PAL [8]	69.37 ± 0.64	84.40 ± 0.44
	inv-equ [39]	67.28 ± 0.80	84.78 ± 0.50
	CSEI [40]	68.94 ± 0.28	85.07 ± 0.50
	COSOC [9]	69.28 ± 0.49	85.16 ± 0.42
	EASY $2 \times \text{ResNet12} \left(\frac{1}{\sqrt{2}} \right)$ (ours)	70.63 ± 0.20	86.28 ± 0.12
$36M$	S2M2R [12]	64.93 ± 0.18	83.18 ± 0.11
	LR + DC [17]	68.55 ± 0.55	82.88 ± 0.42
	EASY $3 \times \text{ResNet12}$ (ours)	71.75 ± 0.19	87.15 ± 0.12

TABLE VI

1-SHOT AND 5-SHOT ACCURACY OF STATE-OF-THE-ART METHODS AND PROPOSED SOLUTION ON **MINIIMAGENET** IN **TRANSDUCTIVE** SETTING.

	Method	1-shot	5-shot
$\leq 12M$	TIM-GD [42]	73.90	85.00
	ODC [43]	77.20 ± 0.36	87.11 ± 0.42
	PEM _n E-BMS* [32]	80.56 ± 0.27	87.98 ± 0.14
	SSR [44]	68.10 ± 0.60	76.90 ± 0.40
	iLPC [45]	69.79 ± 0.99	79.82 ± 0.55
	EPNet [31]	66.50 ± 0.89	81.60 ± 0.60
	DPGN [46]	67.77 ± 0.32	84.60 ± 0.43
	ECKPN [47]	70.48 ± 0.38	85.42 ± 0.46
	Rot+KD+POODLE [48]	77.56	85.81
	EASY $2 \times \text{ResNet12} \left(\frac{1}{\sqrt{2}} \right)$ (ours)	82.31 ± 0.24	88.57 ± 0.12
$36M$	SSR [44]	72.40 ± 0.60	80.20 ± 0.40
	fine-tuning(train+val) [49]	68.11 ± 0.69	80.36 ± 0.50
	SIB+E ³ BM [50]	71.40	81.20
	LR+DC [17]	68.57 ± 0.55	82.88 ± 0.42
	EPNet [31]	70.74 ± 0.85	84.34 ± 0.53
	TIM-GD [42]	77.80	87.40
	PT+MAP [51]	82.92 ± 0.26	88.82 ± 0.13
	iLPC [45]	83.05 ± 0.79	88.82 ± 0.42
	ODC [43]	80.64 ± 0.34	89.39 ± 0.39
	PEM _n E-BMS* [32]	83.35 ± 0.25	89.53 ± 0.13
		EASY $3 \times \text{ResNet12}$ (ours)	84.04 ± 0.23

TAKE-HOME MESSAGES

Deep learning is getting prevalent in society

It has many applications in many fields

It has become state of the art in most domains

- Introduction
- From neural networks to deep learning
- In practice:
 - Transfer learning
 - Standard libraries
- Limitations & future
- Focus: few-shot learning



TAKE-HOME MESSAGES

Various architectures tackle various tasks

Backpropagation is very important

Beware of overfitting!

- Introduction
- From neural networks to deep learning
- In practice:
 - Transfer learning
 - Standard libraries
- Limitations & future
- Focus: few-shot learning



TAKE-HOME MESSAGES

This is what most people do in practice

Allows exploiting already trained architectures

A NN consists of a backbone and a classifier

Training the classifier is necessary

- Introduction
- From neural networks to deep learning
- In practice:
 - Transfer learning
 - Standard libraries
- Limitations & future
- Focus: few-shot learning



TAKE-HOME MESSAGES

Powerful options such as Keras or Torch exist
Do not reinvent the wheel!

- Introduction
- From neural networks to deep learning
- In practice:
 - Transfer learning
 - Standard libraries
- Limitations & future
- Focus: few-shot learning



TAKE-HOME MESSAGES

There is a big shift toward transformers

High ecological impact

Research in the field is very vast

- Introduction
- From neural networks to deep learning
- In practice:
 - Transfer learning
 - Standard libraries
- Limitations & future
- Focus: few-shot learning



TAKE-HOME MESSAGES

Sometimes, we don't have enough data
even for transfer learning


Solutions exist and employ
semi-supervised techniques

There are many more questions linked
to few-shot learning, e.g., data selection


- Introduction
- From neural networks to deep learning
- In practice:
 - Transfer learning
 - Standard libraries
- Limitations & future
- Focus: few-shot learning

Workshop on FSL
IMT – 2-6/12/2024

<https://wfsl.istc-conf.org/>



TAKE-HOME MESSAGES




<https://www.imt-atlantique.fr/en/research-innovation/teams/brain>



- Introduction
- From neural networks to deep learning
- In practice:
 - Transfer learning
 - Standard libraries
- Limitations & future
- Focus: few-shot learning

Workshop on FSL
IMT – 2-6/12/2024

<https://wfsl.istc-conf.org/>



(DEEP) NEURAL NETWORKS

